

KU Leuven Summer School
Segment 1
Jumping into Bayes

Paul Gustafson

September 15, 2022

Example story

Going to be getting some patient-level data:

- ▶ take standard ($X = 0$) or new ($X = 1$) drug;
- ▶ symptoms improve ($Y = 1$) or not ($Y = 0$).

Define $p_x = \Pr(Y = 1|X = x)$, $x = 0, 1$.

In advance of receiving data, experts feels that $p_x < 0.05$ is unlikely, as is $p_x > 0.33$, for both $x = 0$ and $x = 1$.

This expert knowledge is encoded in a **prior distribution**:

$p_0, p_1 \stackrel{\text{iid}}{\sim} \text{beta}(a, b)$.

```
### declare our hyperparameter specification
```

```
hyp <- list(a=3, b=15)
```

```
### and confirm this specification is fit-for-purpose
```

```
qbeta(c(.05, .95), hyp$a, hyp$b)
```

```
## [1] 0.0499 0.3262
```

Arrival of the data

```
dim(dat)
```

```
## [1] 37 2
```

```
head(dat)
```

```
##   x y  
## 1 1 0  
## 2 0 0  
## 3 0 0  
## 4 0 0  
## 5 0 0  
## 6 0 0
```

```
table(dat)
```

```
##      y  
## x | 0 1  
## 0 | 22 1  
## 1 | 9 5
```

	0	1
0	22	1
1	9	5

37

Statistical model for data given parameters

Sufficient statistics $S_j = \sum_{i: x_i=j} y_i$, $j = 0, 1$.

$S_j \sim \text{Bin}(n_j, p_j)$, independently for $j = 0, 1$.

```
n <- as.vector(table(dat$x))  
s.dat <- c(sum(dat$y[dat$x==0]), sum(dat$y[dat$x==1]))
```

```
n
```

```
## [1] 23 14
```

```
s.dat
```

```
## [1] 1 5
```

Scalar parameter of most interest:

$$\psi = \text{logit}(p_1) - \text{logit}(p_0) = \log\left(\frac{p_1}{1-p_1}\right) - \log\left(\frac{p_0}{1-p_0}\right) = \log \text{OR}(X, Y)$$

Generative model

The amalgamation of the prior distribution and the statistical model can be referred to as the *generative model*.

- ▶ A completely specified *joint* dist. for data and parameters.
- ▶ May feel weird, since from the investigator's perspective parameters are **fixed but unknown**, not **random**. Need to remember: prior distribution is being used to describe **investigator knowledge** about these fixed but unknown quantities.
- ▶ Perhaps helpful heuristic: 'Mother Nature' uses the generative model to generate both the state of the world and the ensuing data. But she keeps the former private, only shares the data with the investigator.

true param.
values

Posterior distribution

Describes investigator's knowledge about the state of the world after Nature has shared the data.

$$f(p|s) = \frac{f(s|p)f(p)}{f(s)}$$

proportional
as a function
of p , for
fixed s

$$\propto f(s|p)f(p)$$

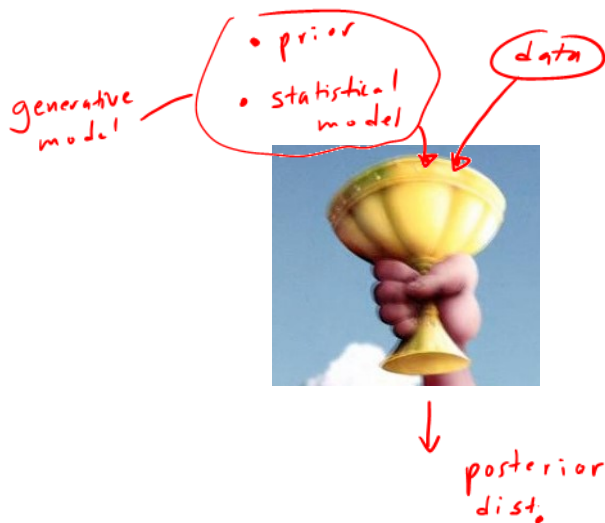
likelihood /
prior

Bayes
theorem

sometimes
see more
explicitly
with

$$f(s) = \int f(s|\tilde{p})f(\tilde{p}) d\tilde{p}$$

The Holy Grail of Bayesian Statistics



Determining the posterior distribution - in a literal, brute-force manner

"Bayesian socks" - video

```
### represent the posterior via m MC realizations
```

```
m <- 20000; opt.brute <- matrix(NA, m, 2)
```

```
colnames(opt.brute) <- c("p0", "p1")
```

```
have <- 0
```

```
while (have < m) {
```

```
### simulate from generative model
```

```
p.0.sim <- rbeta(1, hyp$a, hyp$b)
```

```
p.1.sim <- rbeta(1, hyp$a, hyp$b)
```

```
s.sim <- c(rbinom(1, size=n[1], prob=p.0.sim),  
          rbinom(1, size=n[2], prob=p.1.sim))
```

```
### only keep if the simulated data matches observed
```

```
if (all(s.sim==s.dat)) {
```

```
  have <- have + 1
```


```
  opt.brute[have,] <- c(p.0.sim, p.1.sim)
```

```
}
```



```
}
```


Prior and posterior distributions of (p_0, p_1)

impact of data



Estimate the target



Do some inference

A point estimate, $E(\psi|\text{Data})$:

posterior
mean

```
mean(opt.brute[, "psi"])
```

```
## [1] 1.2
```

A “standard error” $SD(\psi|\text{Data})$:

posterior SD

```
sqrt(var(opt.brute[, "psi"]))
```

```
## [1] 0.705
```

A 95% interval estimate:

```
quantile(opt.brute[, "psi"], c(0.025, 0.975))
```

```
##    2.5%  97.5%
```

```
## -0.111  2.644
```

Nuance: Statistical error versus numerical / Monte Carlo error

$SD(\psi|\text{Data})$ is a 'standard-error-like' quantification of how well the sample quantity $E(\psi|\text{Data})$ estimates the population parameter ψ .

That said, carefully distinguish the roles of:

```
sqrt(var(opt.brute[, "psi"]))
```

```
## [1] 0.705
```

} governed by n

- numerical approx of $SD(\psi|\text{Data})$
- becomes exact as $m \rightarrow \infty$

last slide $\hat{\psi} = \frac{1}{m} \sum_{i=1}^m \psi^{(i)} = 1.20$
posterior mean

```
sqrt(var(opt.brute[, "psi"])/m)
```

```
## [1] 0.00498
```

— so numerical approx good
to about $\pm 2(0.05) = \pm 0.1$
→ Monte Carlo SE

Answer by package - JAGS

JAGS syntax, not R code

```
genmod.JAGS <- "model{  
  
  ### prior distribution  
  p0 ~ dbeta(a,b)  
  p1 ~ dbeta(a,b)  
  
  ### statistical model  
  s0 ~ dbinom(p0, n0)  
  s1 ~ dbinom(p1, n1)  
  
  ### and for convenience, store the target param also  
  psi <- logit(p1)-logit(p0)  
}"
```

encoding
gen.
model

Answer by JAGS, continued

```
require(rjags)
```

```
### generative model, data go in
```

```
mod <- jags.model(textConnection(genmod.JAGS),  
                  data=list(s0=s.dat[1], n0=n[1],  
                             s1=s.dat[2], n1=n[2],  
                             a=hyp$a, b=hyp$b),  
                  n.chains=4)
```

hyperparams

```
### bit of burn-in
```

```
update(mod, 2000)
```

```
### MC output comes out
```

```
opt.JAGS <- coda.samples(mod,  
                           variable.names=c("psi"),  
                           n.iter=10000)
```

Can we find our friendly neighbourhood estimates in the package output?

```
summary(opt.JAGS)
```

Yes, at least to some level of numerical precision

```
##
## Iterations = 3001:13000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
```

Mean	SD	Naive SE	Time-series SE
1.19930	0.69515	0.00348	0.00450

```
## 2. Quantiles for each variable:
```

2.5%	25%	50%	75%	97.5%
-0.105	0.724	1.179	1.649	2.625

||
reflecting numerical (MC) precision

FYI, lots of output options

```
require(MCMCvis)
MCMCsummary(opt.JAGS)
```

##	mean	sd	2.5%	50%	97.5%	Rhat	n.eff
## psi	1.2	0.695	-0.105	1.18	2.63	1	23849

different
quantification of
numerical precision
as good as if
we had this

many
independent

MC
draws

we asked for 4×10^6 ,
but they aren't independent

ASIDE: off-the-shelf **non-Bayesian** analysis of these data?

```
### ML estimate of psi
```

```
logit(s.dat[2]/n[2]) - logit(s.dat[1]/n[1])
```

```
## [1] 2.5
```

```
### standard error
```

```
sqrt(sum(1/c(s.dat[1], n[1]-s.dat[1], s.dat[2], n[2]-s.dat[2])))
```

```
## [1] 1.16
```

not close to
Bayes answers

```
### or if you prefer, get these from...
```

```
glm(y~x, family=binomial)
```

Can speculate on **two** reasons why our present Bayesian answer is not so close to this

- no prior
- sparse data

Black-box package actually overkill for problem we just did

Our problem has **conjugate** structure, ergo a clean math description of the posterior dist. for $(p_0, p_1 | S_0, S_1)$:

$(p_j | S = s) \sim \text{beta}(a + s_j, b + (n_j - s_j))$, independently for $j = 0, 1$.

So yet another route to our inference

```
opt.MC <- cbind(  
  "p0" = rbeta(m, hyp$a+s.dat[1], hyp$b+n[1]-s.dat[1]),  
  "p1" = rbeta(m, hyp$a+s.dat[2], hyp$b+n[2]-s.dat[2]))  
  
opt.MC <- cbind(opt.MC,  
  "psi" = logit(opt.MC[, "p1"]) - logit(opt.MC[, "p0"]))
```

```
## posterior mean and SD of target  
c(mean(opt.MC[, "psi"]), sqrt(var(opt.MC[, "psi"])))
```

```
## [1] 1.196 0.697
```

✓ *same answer*

```
### quality of numerical approximation  
sqrt(var(opt.MC[, "psi"])/m)
```

```
## [1] 0.00493
```

In fact, for problem at hand, can skip MC approximation,
compute exactly

```
### posterior mean of psi
```

```
(digamma(hyp$a+s.dat[2])-digamma(hyp$b+n[2]-s.dat[2])) -  
(digamma(hyp$a+s.dat[1]) - digamma(hyp$b+n[1]-s.dat[1]))
```

```
## [1] 1.2
```



same

```
### posterior SD of psi
```

```
sqrt(trigamma(hyp$a+s.dat[2]) + trigamma(hyp$b+n[2]-s.dat[2]) +  
      trigamma(hyp$a+s.dat[1]) + trigamma(hyp$b+n[1]-s.dat[1]))
```

```
## [1] 0.698
```



Back to the holy grail

What *might* a platinum grail spit out?

exact values of any posterior quantity

What *might* a golden grail spit out?

MC realizations $\theta^{(1)}, \dots, \theta^{(m)}$ which are iid

What *might* a silver grail spit out?

$\theta^{(1)}, \dots, \theta^{(m)}$ are dependent

from the
posterior

What **does** a bronze grail spit out?

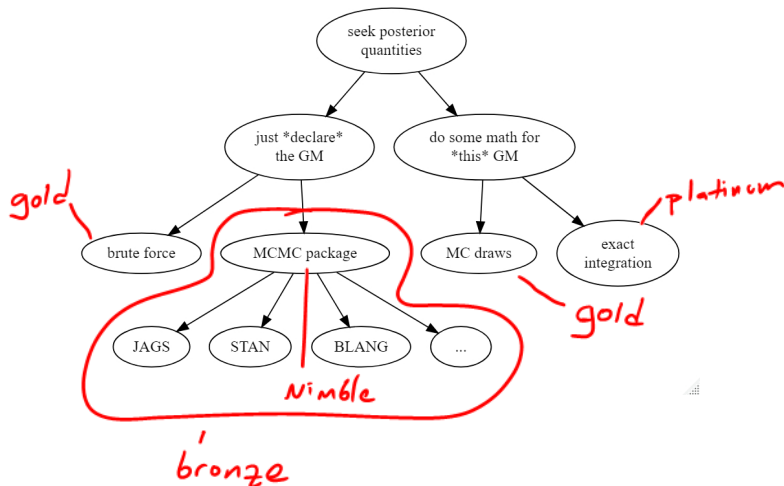
$f(\theta | \text{data})$

$\theta^{(1)}, \dots, \theta^{(m)}$

$\text{dist } \theta^{(m)} \xrightarrow{\text{converges quickly}} f(\theta | \text{data}) \text{ as } m \rightarrow \infty$

$\theta^{(1)} \rightarrow \theta^{(2)} \rightarrow \theta^{(3)} \dots$ Markov chain
↑ arbitrary
clever specification of
the $\theta^{(i)} \rightarrow \theta^{(i+1)}$ recipe

Taxonomy of the computing options we've seen today



Since we are typically stuck with a bronze grail (i.e., MCMC) ...

choice of m
- how many chains
- how long is each

- ▶ Users can/should exert control over how much computing to do.
- ▶ Users can/should monitor how successful this computing is.

- Monte Carlo SE
- m_{eff}
- R_{hat}
- traceplots
- ...

APPENDIX: Could you stand to see one more computational route to the same end?

```
genmod.STAN <- "  
data {  
  int<lower=0> n[2];  
  int<lower=0> s[2];  
  real<lower=0> a[2];  
  real<lower=0> b[2];  
}  
parameters {  
  real<lower=0,upper=1> p[2];  
}  
transformed parameters {  
  real psi;  
  psi =logit(p[2])-logit(p[1]);  
}  
model {  
  p ~ beta(a,b);  
  s ~ binomial(n,p);  
}
```


STAN, continued

```
require("rstan"); rstan_options(auto_write = TRUE)
```

```
opt.STAN <- stan(model_code=genmod.STAN,  
  data=list(s=s.dat, n=n, a=rep(hyp$a,2), b=rep(hyp$b,2)),  
  iter=12000)
```

STAN, continued

```
summary(opt.STAN)$summary
```

##	mean	se_mean	sd	2.5%	25%	50%	75%
## p[1]	0.0965	0.000304	0.0447	0.0286	0.0631	0.0902	0.123
## p[2]	0.2496	0.000529	0.0746	0.1200	0.1961	0.2442	0.298
## psi	1.2074	0.004873	0.6817	-0.0863	0.7423	1.1863	1.654
## lp__	-32.1083	0.009706	1.0006	-34.7985	-32.4934	-31.8057	-31.400
##	97.5%	n_eff	Rhat				
## p[1]	0.199	21618	1				
## p[2]	0.405	19883	1				
## psi	2.601	19572	1				
## lp__	-31.127	10627	1				

Due diligence: Two different black-boxes give same answer?

```
MCMCsummary(opt.STAN)
```

##	mean	sd	2.5%	50%	97.5%	Rhat	n.eff
## p[1]	0.0965	0.0447	0.0286	0.0902	0.199	1	21618
## p[2]	0.2496	0.0746	0.1200	0.2442	0.405	1	19883
## psi	1.2074	0.6817	-0.0863	1.1863	2.601	1	19572
## lp__	-32.1083	1.0006	-34.7985	-31.8057	-31.127	1	10627

```
MCMCsummary(opt.JAGS)
```

##	mean	sd	2.5%	50%	97.5%	Rhat	n.eff
## psi	1.2	0.695	-0.105	1.18	2.63	1	23849